

```
import pandas as pd
df=pd.read_csv(r"/content/laptop_price.csv")
print(df.head())
```

	Company	Product	TypeName	Inches	ScreenResolution \
0	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600
1	Apple	Macbook Air	Ultrabook	13.3	1440x900
2	HP	250 G6	Notebook	15.6	Full HD 1920x1080
3	Apple	MacBook Pro	Ultrabook	15.4	IPS Panel Retina Display 2880x1800
4	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600

	CPU_Company	CPU_Type	CPU_Frequency (GHz)	RAM (GB)	\
0	Intel	Core i5	2.3	8	
1	Intel	Core i5	1.8	8	
2	Intel	Core i5 7200U	2.5	8	
3	Intel	Core i7	2.7	16	
4	Intel	Core i5	3.1	8	

	Memory	GPU_Company	GPU_Type	OpSys \
0	128GB SSD	Intel	Iris Plus Graphics 640	macOS
1	128GB Flash Storage	Intel	HD Graphics 6000	macOS
2	256GB SSD	Intel	HD Graphics 620	No OS
3	512GB SSD	AMD	Radeon Pro 455	macOS
4	256GB SSD	Intel	Iris Plus Graphics 650	macOS

	Weight (kg)	Price (Euro)
0	1.37	1339.69
1	1.34	898.94
2	1.86	575.00
3	1.83	2537.45
4	1.37	1803.60

```
# Check the first few rows
print(df.head())

# Get basic information about the dataset
print(df.info())

# Summary statistics for numerical columns
print(df.describe())

# Check for missing values
print(df.isnull().sum())
```

	Company	Product	TypeName	Inches	ScreenResolution \
0	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600
1	Apple	Macbook Air	Ultrabook	13.3	1440x900
2	HP	250 G6	Notebook	15.6	Full HD 1920x1080
3	Apple	MacBook Pro	Ultrabook	15.4	IPS Panel Retina Display 2880x1800
4	Apple	MacBook Pro	Ultrabook	13.3	IPS Panel Retina Display 2560x1600

	CPU_Company	CPU_Type	CPU_Frequency (GHz)	RAM (GB)	\
0	Intel	Core i5	2.3	8	
1	Intel	Core i5	1.8	8	
2	Intel	Core i5 7200U	2.5	8	
3	Intel	Core i7	2.7	16	
4	Intel	Core i5	3.1	8	

	Memory	GPU_Company	GPU_Type	OpSys \
0	128GB SSD	Intel	Iris Plus Graphics 640	macOS
1	128GB Flash Storage	Intel	HD Graphics 6000	macOS
2	256GB SSD	Intel	HD Graphics 620	No OS
3	512GB SSD	AMD	Radeon Pro 455	macOS
4	256GB SSD	Intel	Iris Plus Graphics 650	macOS

	Weight (kg)	Price (Euro)
0	1.37	1339.69
1	1.34	898.94
2	1.86	575.00
3	1.83	2537.45
4	1.37	1803.60

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1275 entries, 0 to 1274

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	Company	1275 non-null	object
1	Product	1275 non-null	object
2	TypeName	1275 non-null	object
3	Inches	1275 non-null	float64
4	ScreenResolution	1275 non-null	object
5	CPU_Company	1275 non-null	object
6	CPU_Type	1275 non-null	object
7	CPU_Frequency (GHz)	1275 non-null	float64
8	RAM (GB)	1275 non-null	int64
9	Memory	1275 non-null	object
10	GPU_Company	1275 non-null	object
11	GPU_Type	1275 non-null	object

```

12 OpSys          1275 non-null object
13 Weight (kg)    1275 non-null float64
14 Price (Euro)   1275 non-null float64

```

```
dtypes: float64(4), int64(1), object(10)
```

```
memory usage: 149.5+ KB
```

```
None
```

	Inches	CPU_Frequency (GHz)	RAM (GB)	Weight (kg)	\
count	1275.000000	1275.000000	1275.000000	1275.000000	
mean	15.022902	2.302980	8.440784	2.040525	
std	1.429470	0.503846	5.097809	0.669196	
min	10.100000	0.900000	2.000000	0.690000	
25%	14.000000	2.000000	4.000000	1.500000	
50%	15.600000	2.500000	8.000000	2.040000	
75%	15.600000	2.700000	8.000000	2.310000	
max	18.400000	3.600000	64.000000	4.700000	

	Price (Euro)
count	1275.000000
mean	1134.969059
std	700.752504
min	174.000000
25%	609.000000
50%	989.000000
75%	1496.500000
max	6099.000000

```

Company          0
Product          0
TypeName         0
Inches           0
ScreenResolution 0
CPU_Company      0
CPU_Type         0
CPU_Frequency (GHz) 0
RAM (GB)         0
Memory           0
GPU_Company      0
GPU_Type         0
OpSys            0
Weight (kg)      0
Price (Euro)     0
dtype: int64

```

```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

```

```
# Load the dataset
```

```
file_path = 'laptop_price.csv' # Replace with the correct path to
your file
```

```
df = pd.read_csv(file_path)
```

```

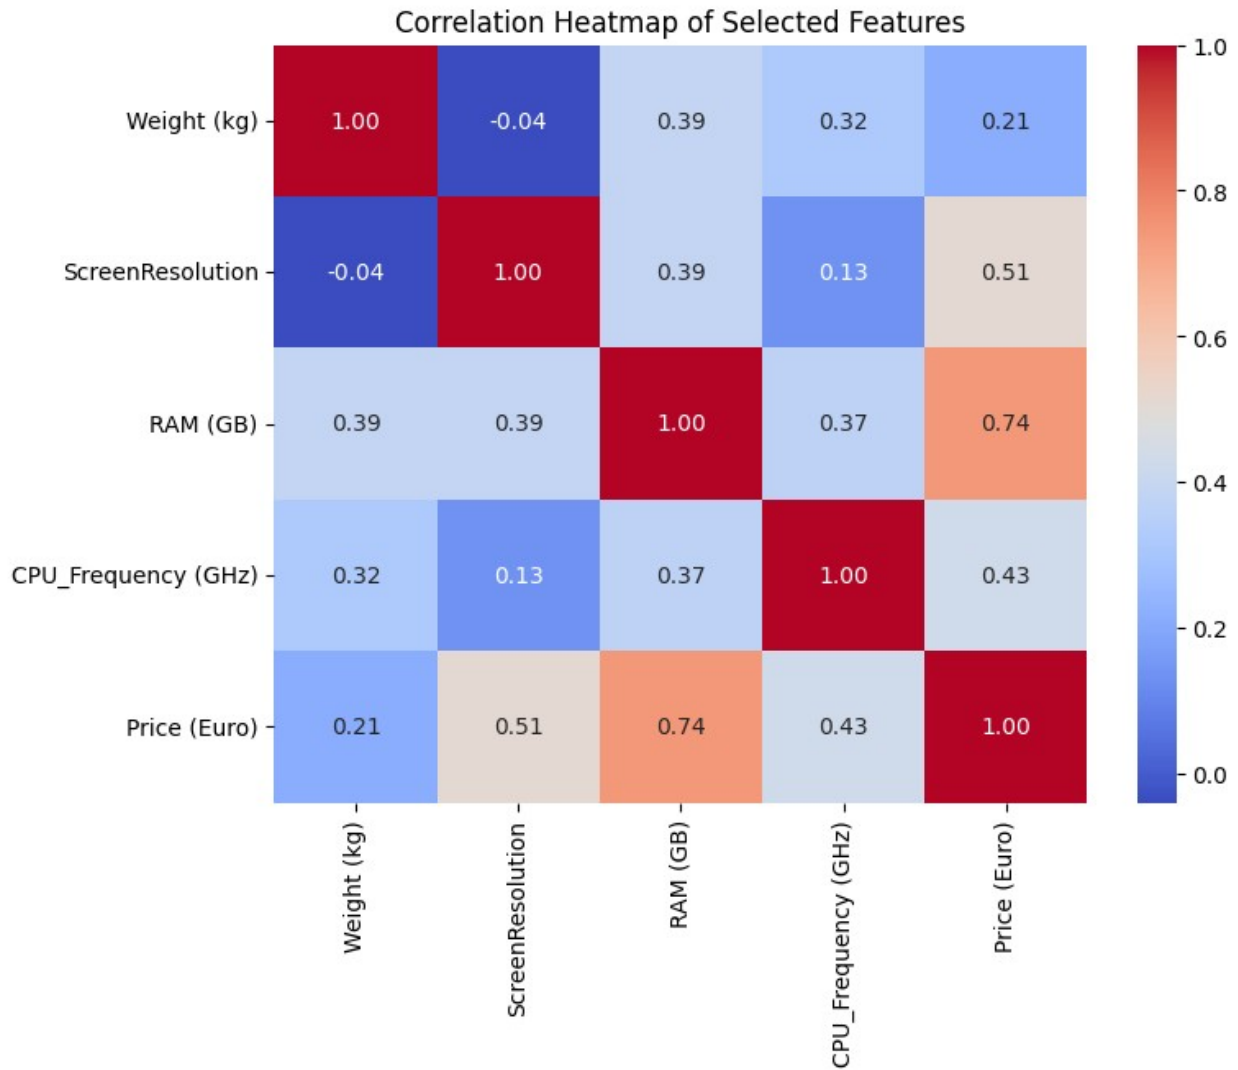
# Derive 'ScreenResolution' and 'Memory_Size' features if needed
df['ScreenResolution'] = df['ScreenResolution'].str.extract(r'(\d+)x(\d+)').astype(float).prod(axis=1) # Pixel count approximation
df['Memory_Size (GB)'] = df['Memory'].str.extract(r'(\d+)').astype(int) # Extract memory size in GB

# Selecting relevant features
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)", "CPU_Frequency (GHz)", "Price (Euro)"]
df_selected = df[selected_features]

# Correlation matrix
correlation_matrix = df_selected.corr()

# Plotting the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm",
fmt=".2f", cbar=True)
plt.title("Correlation Heatmap of Selected Features")
plt.show()

```



```
import pandas as pd

# Load the dataset
file_path = 'laptop_price.csv' # Replace with the correct path to your file
df = pd.read_csv(file_path)

# Derive 'ScreenResolution' and 'Memory_Size' features if needed
df['ScreenResolution'] = df['ScreenResolution'].str.extract(r'(\d+)\s(\d+)').astype(float).prod(axis=1) # Pixel count approximation
df['Memory_Size (GB)'] = df['Memory'].str.extract(r'(\d+)\sGB').astype(int) # Extract memory size in GB

# Selecting relevant numerical features
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)", "CPU_Frequency (GHz)", "Price (Euro)"]
```

```

df_selected = df[selected_features]

# Compute the correlation matrix
correlation_matrix = df_selected.corr()

# Display the correlation matrix
print("Correlation Matrix (Numerical):")
print(correlation_matrix)

Correlation Matrix (Numerical):

```

	Weight (kg)	ScreenResolution	RAM (GB)	\
Weight (kg)	1.000000	-0.040126	0.389370	
ScreenResolution	-0.040126	1.000000	0.388002	
RAM (GB)	0.389370	0.388002	1.000000	
CPU_Frequency (GHz)	0.318649	0.130565	0.366254	
Price (Euro)	0.211883	0.511753	0.740287	

	CPU_Frequency (GHz)	Price (Euro)
Weight (kg)	0.318649	0.211883
ScreenResolution	0.130565	0.511753
RAM (GB)	0.366254	0.740287
CPU_Frequency (GHz)	1.000000	0.428847
Price (Euro)	0.428847	1.000000

```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load dataset
df = pd.read_csv("laptop_price.csv")

# Step 1: Convert 'Price (Euro)' to binary classification
threshold = df['Price (Euro)'].median() # Define the threshold for
classification
df['Price_Category'] = (df['Price (Euro)'] > threshold).astype(int) #
1 if price > threshold, 0 if price <= threshold

# Step 2: Selected features
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
# Adjust as needed

# Step 3: Pair plot
print("Pair Plot:")
sns.pairplot(df[selected_features + ["Price_Category"]],
hue="Price_Category", palette="coolwarm")
plt.show()

# Step 4: Histograms
print("Histograms:")
df[selected_features].hist(bins=15, figsize=(10, 8), color='skyblue',

```

```

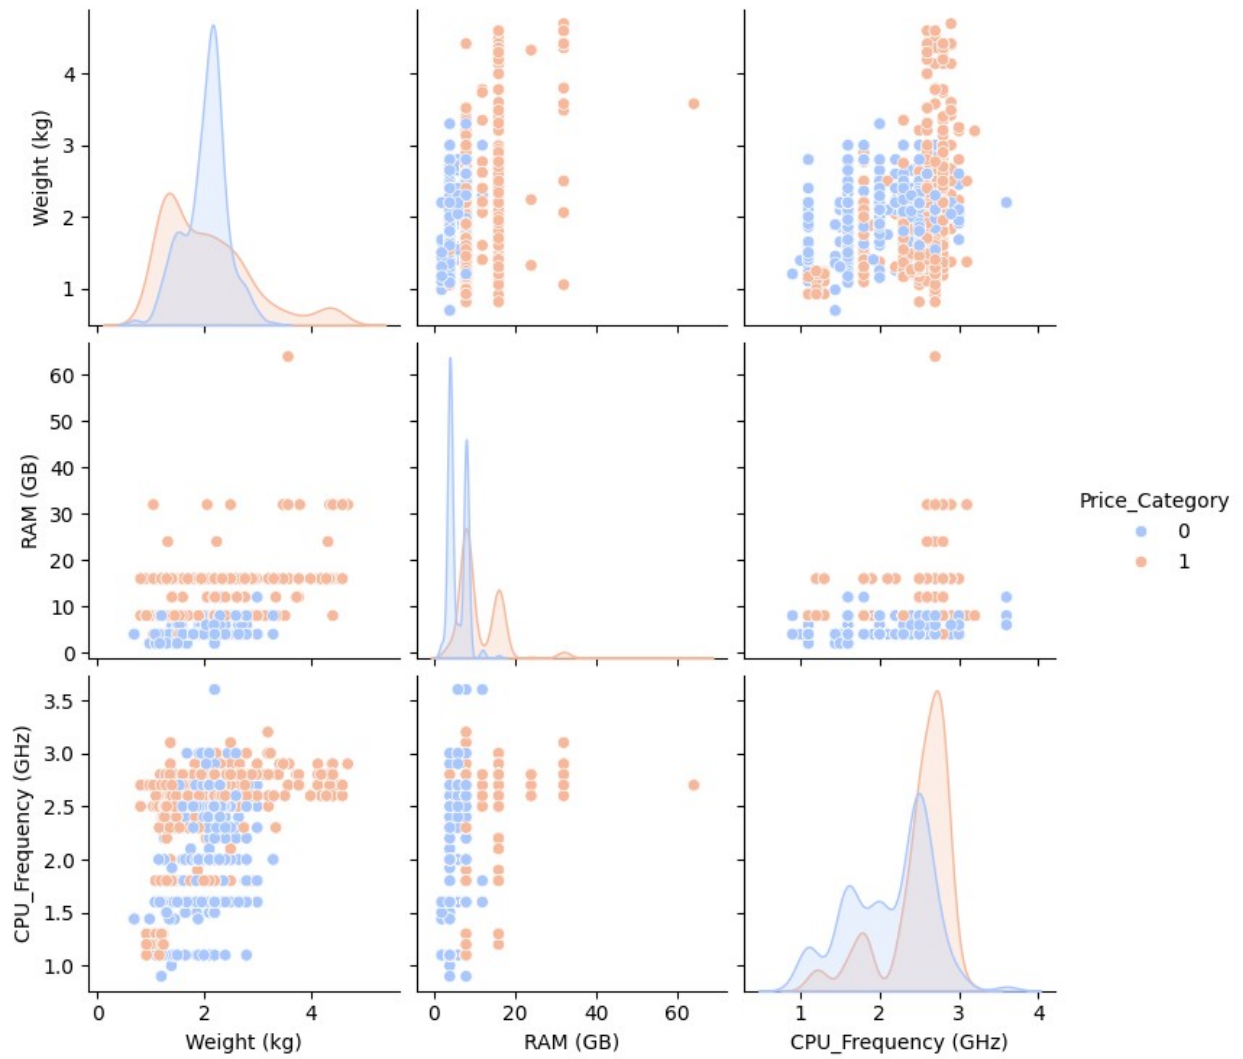
edgecolor='black')
plt.suptitle('Feature Distributions', size=16)
plt.show()

# Step 5: Correlation Heatmap
print("Correlation Heatmap:")
corr = df[selected_features + ["Price_Category"]].corr()
plt.figure(figsize=(8, 6))
sns.heatmap(corr, annot=True, cmap="coolwarm", fmt=".2f",
linewidths=0.5)
plt.title("Correlation Heatmap")
plt.show()

# Step 6: Scatter Plot (Weight vs CPU_Frequency, colored by
Price_Category)
print("Scatter Plot:")
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x="Weight (kg)", y="CPU_Frequency (GHz)",
hue="Price_Category", palette="coolwarm")
plt.title("Weight vs CPU Frequency (Colored by Price Category)")
plt.xlabel("Weight (kg)")
plt.ylabel("CPU Frequency (GHz)")
plt.show()

```


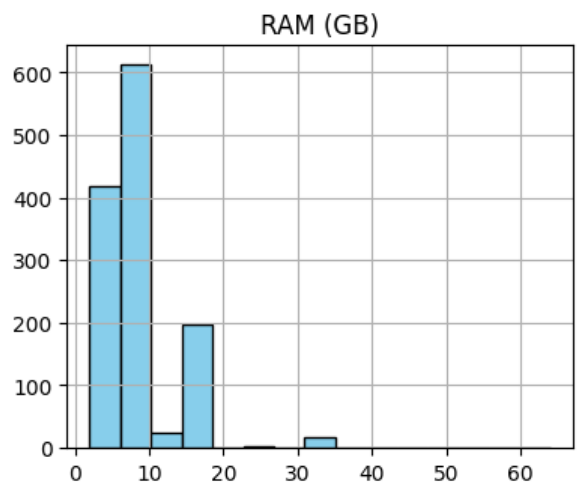

Pair Plot:




Histograms:




## Feature Distributions



Correlation Heatmap:



Scatter Plot:



```
import pandas as pd
```

```
# Derive 'ScreenResolution' and 'Memory_Size' features if needed
df['ScreenResolution'] = df['ScreenResolution'].str.extract(r'(\d+)x(\d+)').astype(float).prod(axis=1) # Pixel count approximation
df['Memory_Size (GB)'] = df['Memory'].str.extract(r'(\d+)').astype(int) # Extract memory size in GB
```

```
# Define the selected features and target
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)", "CPU_Frequency (GHz)"] # Based on correlation
X = df[selected_features] # Features
y = df['Price (Euro)'] # Target
```

```
# Display the first few rows of X and y
print("Selected Features (X):")
print(X.head())
```

```
print("\nTarget (y):")
print(y.head())
```

Selected Features (X):

	Weight (kg)	ScreenResolution	RAM (GB)	CPU_Frequency (GHz)
0	1.37	4096000.0	8	2.3
1	1.34	1296000.0	8	1.8
2	1.86	2073600.0	8	2.5
3	1.83	5184000.0	16	2.7
4	1.37	4096000.0	8	3.1

Target (y):

0	1339.69
1	898.94
2	575.00
3	2537.45
4	1803.60

Name: Price (Euro), dtype: float64

```
import pandas as pd
import numpy as np
```

*# Load the dataset*

*file\_path = 'laptop\_price.csv' # Replace with the correct path to your file*

```
df = pd.read_csv(file_path)
```

*# Step 1: Check for missing values*

```
print("Missing Values:\n", df.isnull().sum())
```

*# Step 2: Handle missing values*

*# Assuming no missing values in key columns; fill any missing numeric data with the mean*

```
df.fillna(df.mean(numeric_only=True), inplace=True)
```

*# Step 3: Convert categorical data to numeric*

*# Example: Encoding the 'OpSys' column using one-hot encoding*

```
if 'OpSys' in df.columns:
```

```
    df = pd.get_dummies(df, columns=['OpSys'], prefix='OpSys',
drop_first=True)
```

*# Step 4: Extract useful information from text columns*

*# Example 1: Extract screen resolution (if not already done)*

```
if 'ScreenResolution' in df.columns:
```

```
    resolution_split = df['ScreenResolution'].str.extract(r'(\d+)\s{x}\s(\d+)\s')
```

```
    df['ScreenResolution'] = resolution_split[0].astype(float) *
resolution_split[1].astype(float)
```

*# Example 2: Extract numeric memory size from 'Memory' column*

```

if 'Memory' in df.columns:
    df['Memory_Size (GB)'] = df['Memory'].str.extract(r'(\d+)').astype(float)

# Step 5: Normalize or scale numerical columns (optional, depending on
the ML model)
from sklearn.preprocessing import StandardScaler

# List of numerical columns to scale
numerical_columns = ['Weight (kg)', 'ScreenResolution', 'RAM (GB)',
'CPU_Frequency (GHz)', 'Price (Euro)']
scaler = StandardScaler()
df[numerical_columns] = scaler.fit_transform(df[numerical_columns])

# Step 6: Drop redundant or irrelevant columns
# Example: Dropping the original 'Memory' column after extracting
numeric data
columns_to_drop = ['Memory'] # Add any other redundant columns here
df.drop(columns=columns_to_drop, inplace=True, errors='ignore')

# Step 7: Check the cleaned dataset
print("Cleaned DataFrame:")
print(df.head())

# Optional: Save the cleaned dataset
df.to_csv('laptop_price_cleaned.csv', index=False)

```

Missing Values:

Company	0
Product	0
TypeName	0
Inches	0
ScreenResolution	0
CPU_Company	0
CPU_Type	0
CPU_Frequency (GHz)	0
RAM (GB)	0
Memory	0
GPU_Company	0
GPU_Type	0
OpSys	0
Weight (kg)	0
Price (Euro)	0

dtype: int64

Cleaned DataFrame:

	Company	Product	TypeName	Inches	ScreenResolution
0	Apple	MacBook Pro	Ultrabook	13.3	1.376413
1	Apple	Macbook Air	Ultrabook	13.3	-0.634604

```

Intel
2      HP      250 G6   Notebook   15.6      -0.076116
Intel
3      Apple  MacBook Pro Ultrabook   15.4      2.157837
Intel
4      Apple  MacBook Pro Ultrabook   13.3      1.376413
Intel

      CPU_Type  CPU_Frequency (GHz)  RAM (GB) GPU_Company  ... Price
(Euro)  \
0      Core i5          -0.005918 -0.086499      Intel  ...
0.292259
1      Core i5          -0.998674 -0.086499      Intel  ...  -
0.336954
2      Core i5 7200U      0.391185 -0.086499      Intel  ...  -
0.799410
3      Core i7          0.788288  1.483418      AMD    ...
2.002178
4      Core i5          1.582493 -0.086499      Intel  ...
0.954536

      OpSys_Chrome OS  OpSys_Linux  OpSys_Mac OS X  OpSys_No OS  \
0      False          False          False          False
1      False          False          False          False
2      False          False          False          True
3      False          False          False          False
4      False          False          False          False

      OpSys_Windows 10  OpSys_Windows 10 S  OpSys_Windows 7  OpSys_macOS
\
0      False          False          False          True
1      False          False          False          True
2      False          False          False          False
3      False          False          False          True
4      False          False          False          True

      Memory_Size (GB)
0      128.0
1      128.0
2      256.0
3      512.0
4      256.0

[5 rows x 22 columns]

```

```

from sklearn.model_selection import train_test_split

# Define the features (X) and target variable (y)
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Based on analysis
X = df[selected_features] # Features
y = df['Price (Euro)'] # Target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Display the sizes of the splits
print(f"Training Set: {X_train.shape[0]} samples")
print(f"Testing Set: {X_test.shape[0]} samples")

Training Set: 1020 samples
Testing Set: 255 samples

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Step 1: Convert target variable 'Price (Euro)' to binary
classification (e.g., high price vs low price)
# Define a threshold for binary classification (e.g., median price)
threshold = df['Price (Euro)'].median()
y_binary = (y > threshold).astype(int) # 1 if price > threshold, 0 if
price <= threshold

# Step 2: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y_binary,
test_size=0.2, random_state=42)

# Step 3: Standardize the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 4: Initialize the Logistic Regression model
log_reg = LogisticRegression(random_state=42)

# Step 5: Train the model
log_reg.fit(X_train_scaled, y_train)

# Step 6: Predict on the test data
y_pred = log_reg.predict(X_test_scaled)

# Step 7: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)

```

```
print(f"Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Accuracy: 0.7961

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.72	0.77	118
1	0.78	0.86	0.82	137
accuracy			0.80	255
macro avg	0.80	0.79	0.79	255
weighted avg	0.80	0.80	0.79	255

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report,
roc_curve, auc, precision_recall_curve
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

# Assuming 'df' is already loaded, cleaned, and preprocessed
# Define features and target
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
# Example features
X = df[selected_features]
y = (df['Price (Euro)'] > df['Price (Euro)'].median()).astype(int) #
Binary classification

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Train Logistic Regression model
log_reg = LogisticRegression(max_iter=1000, random_state=42)
log_reg.fit(X_train, y_train)

# Predictions
y_pred = log_reg.predict(X_test)
y_pred_proba = log_reg.predict_proba(X_test)[:, 1]

# Step 1: Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
xticklabels=["Below Median", "Above Median"],
yticklabels=["Below Median", "Above Median"])
plt.title("Confusion Matrix")
```



```


plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()


# Step 2: ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--')
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid()
plt.show()


# Step 3: Precision-Recall Curve
precision, recall, thresholds = precision_recall_curve(y_test,
y_pred_proba)
plt.figure(figsize=(8, 6))
plt.plot(recall, precision, color='green', label="Precision-Recall Curve")
plt.title("Precision-Recall Curve")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.legend()
plt.grid()
plt.show()

# Step 4: Feature Importance
coefficients = pd.DataFrame(log_reg.coef_[0], index=selected_features,
columns=['Importance']).sort_values(by='Importance')
plt.figure(figsize=(6, 4))
sns.barplot(x=coefficients['Importance'], y=coefficients.index,
palette="coolwarm")
plt.title("Feature Importance (Logistic Regression Coefficients)")
plt.xlabel("Coefficient Value")
plt.ylabel("Feature")
plt.show()

```



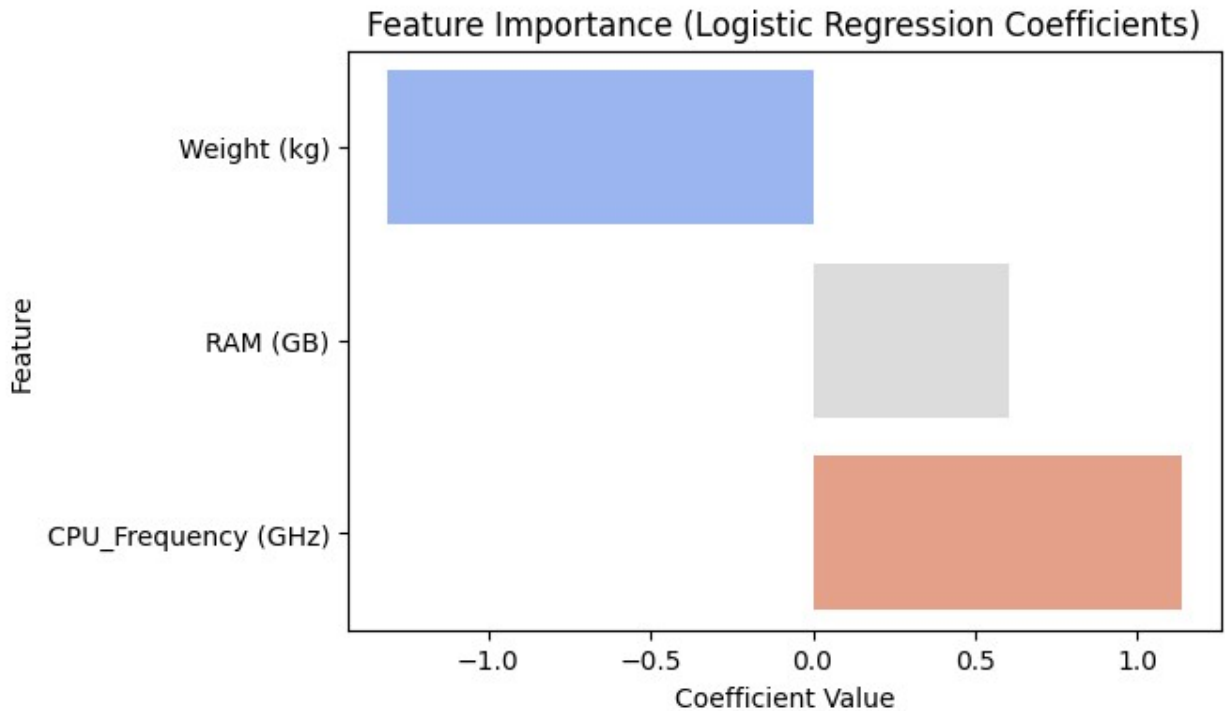




```
<ipython-input-6-ec32ee96fb35>:61: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.barplot(x=coefficients['Importance'], y=coefficients.index, palette="coolwarm")
```



```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Step 1: Convert target variable 'Price (Euro)' to binary
# classification (e.g., high price vs low price)
# Define a threshold for binary classification (e.g., median price)
threshold = df['Price (Euro)'].median()
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 4: Standardize the features (optional)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Step 5: Initialize the Decision Tree Classifier  
dt_classifier = DecisionTreeClassifier(random_state=42)
```

```
# Step 6: Train the model  
dt_classifier.fit(X_train_scaled, y_train)
```

```
# Step 7: Make predictions on the test set  
y_pred = dt_classifier.predict(X_test_scaled)
```

```
# Step 8: Evaluate the model  
accuracy = accuracy_score(y_test, y_pred)  
print(f"Accuracy: {accuracy:.4f}")  
print("Classification Report:")  
print(classification_report(y_test, y_pred))
```

Accuracy: 0.8471

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.86	0.84	118
1	0.88	0.83	0.85	137

accuracy			0.85	255
macro avg	0.85	0.85	0.85	255
weighted avg	0.85	0.85	0.85	255

```
from sklearn.tree import DecisionTreeClassifier, plot_tree  
from sklearn.metrics import confusion_matrix, classification_report  
import seaborn as sns  
import matplotlib.pyplot as plt  
from sklearn.model_selection import train_test_split
```

```
# Define features and target  
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]  
# Example features  
X = df[selected_features]  
y = (df['Price (Euro)'] > df['Price (Euro)'].median()).astype(int) #  
Binary classification
```

```
# Split data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.3, random_state=42)
```

```
# Train Decision Tree model  
tree_clf = DecisionTreeClassifier(max_depth=3, random_state=42) #  
Adjust max_depth as needed  
tree_clf.fit(X_train, y_train)
```

```
# Step 1: Decision Tree Structure Visualization  
plt.figure(figsize=(12, 8))
```

```

plot_tree(tree_clf, feature_names=selected_features,
class_names=["Below Median", "Above Median"],
            filled=True, rounded=True, fontsize=10)
plt.title("Decision Tree Structure")
plt.show()

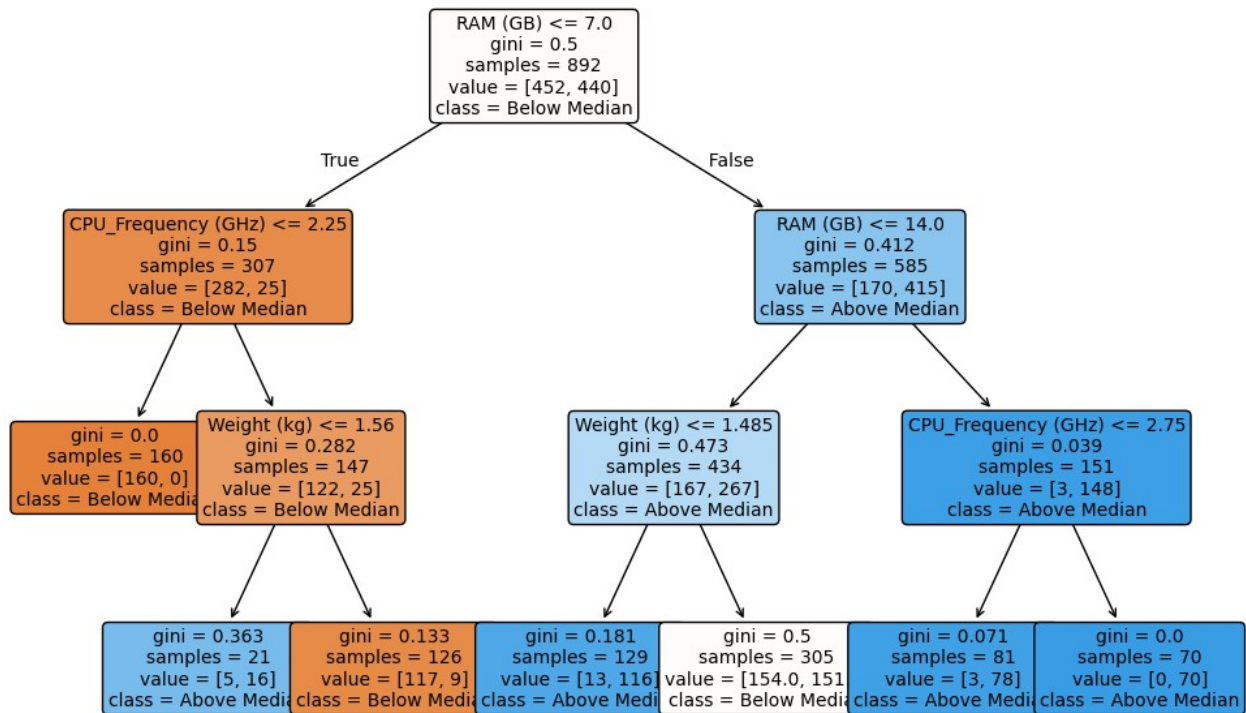
# Step 2: Feature Importance
feature_importances = pd.DataFrame(tree_clf.feature_importances_,
index=selected_features,

columns=["Importance"]).sort_values(by="Importance", ascending=False)
plt.figure(figsize=(6, 4))
sns.barplot(x=feature_importances["Importance"],
y=feature_importances.index, palette="coolwarm")
plt.title("Feature Importance (Decision Tree)")
plt.xlabel("Importance")
plt.ylabel("Feature")
plt.show()

# Step 3: Confusion Matrix
y_pred = tree_clf.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
xticklabels=["Below Median", "Above Median"],
            yticklabels=["Below Median", "Above Median"])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

Decision Tree Structure

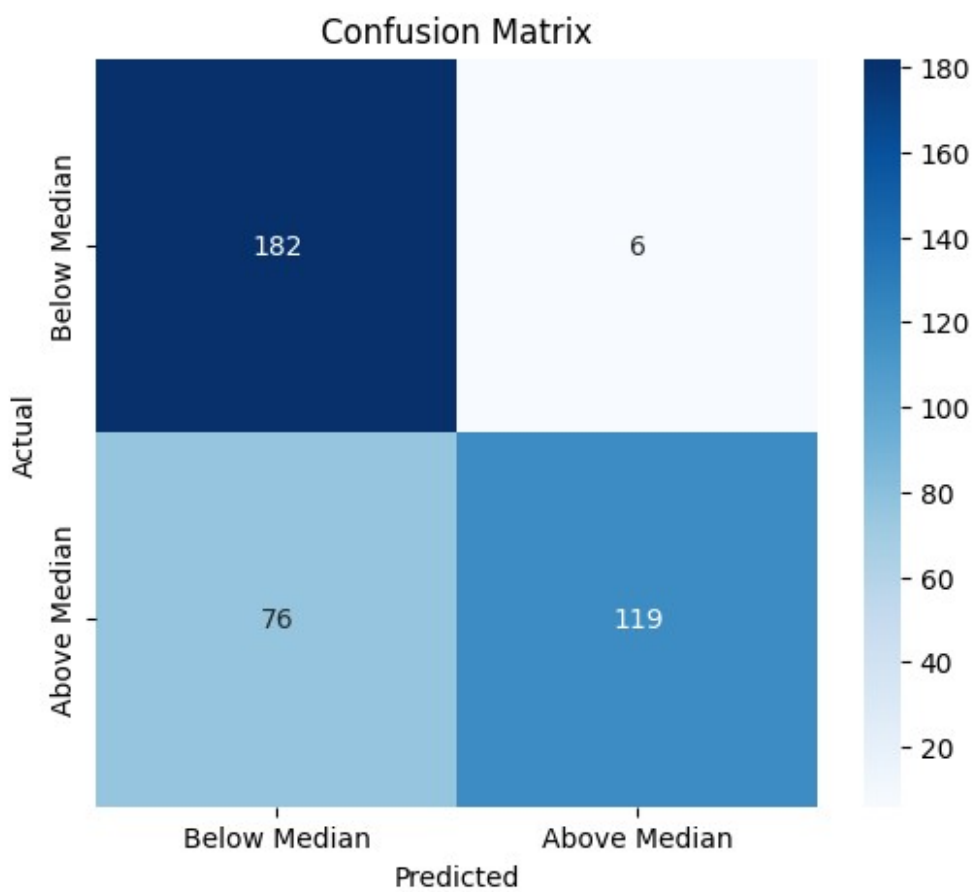
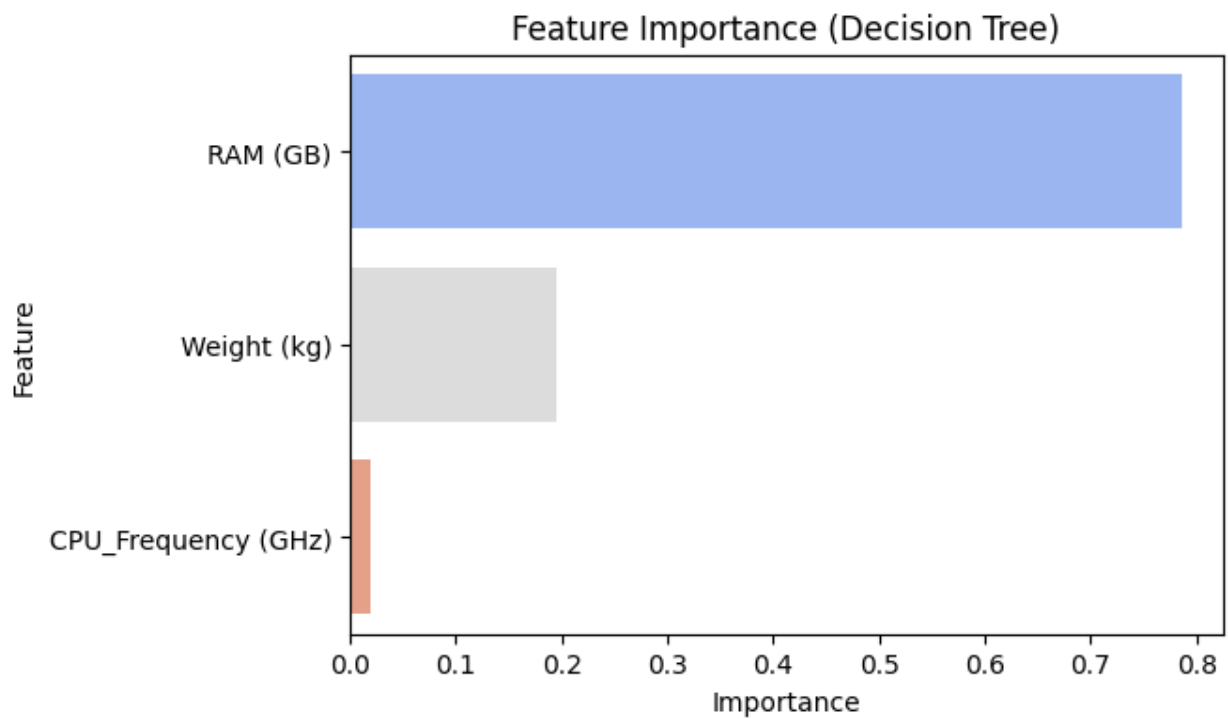


<ipython-input-7-174f91e135b3>:30: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=feature_importances["Importance"],
y=feature_importances.index, palette="coolwarm")
```





```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import pandas as pd

# Step 1: Convert target variable 'Price (Euro)' to binary
# classification (e.g., high price vs low price)
# Define a threshold for binary classification (e.g., median price)
threshold = df['Price (Euro)'].median()
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features and handle the categorical
'ScreenResolution' feature
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Convert 'ScreenResolution' (categorical) to numerical format
using LabelEncoder
label_encoder = LabelEncoder()
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 5: Standardize the features (optional)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 6: Initialize the Random Forest Classifier
rf_classifier = RandomForestClassifier(n_estimators=100,
random_state=42)

# Step 7: Train the model
rf_classifier.fit(X_train_scaled, y_train)

# Step 8: Make predictions on the test set
y_pred = rf_classifier.predict(X_test_scaled)

# Step 9: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

```
<ipython-input-23-69a697354206>:19: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X['ScreenResolution'] =  
label_encoder.fit_transform(X['ScreenResolution'].astype(str))
```

Accuracy: 0.8392

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.83	0.83	118
1	0.85	0.85	0.85	137
accuracy			0.84	255
macro avg	0.84	0.84	0.84	255
weighted avg	0.84	0.84	0.84	255

```
from sklearn.svm import SVC  
from sklearn.metrics import accuracy_score, classification_report  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
  
# Step 1: Convert target variable 'Price (Euro)' to binary  
# classification (e.g., high price vs low price)  
# Define a threshold for binary classification (e.g., median price)  
threshold = df['Price (Euro)'].median()  
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price  
> threshold, 0 if price <= threshold  
  
# Step 2: Select features  
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",  
"CPU_Frequency (GHz)"] # Features  
X = df[selected_features] # Features  
y = y_binary # Target (binary classification)  
  
# Step 3: Split the data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)  
  
# Step 4: Standardize the features (SVM models are sensitive to  
# feature scaling)  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)  
  
# Step 5: Initialize the SVM model
```

```
svm_classifier = SVC(kernel='linear', random_state=42) # You can use
'linear', 'rbf', or 'poly'
```

```
# Step 6: Train the model
```

```
svm_classifier.fit(X_train_scaled, y_train)
```

```
# Step 7: Make predictions on the test set
```

```
y_pred = svm_classifier.predict(X_test_scaled)
```

```
# Step 8: Evaluate the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy:.4f}")
```

```
print("Classification Report:")
```

```
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.7882
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.82	0.69	0.75	118
1	0.77	0.87	0.82	137

accuracy			0.79	255
macro avg	0.79	0.78	0.78	255
weighted avg	0.79	0.79	0.79	255

```
from sklearn.svm import SVC
```

```
from sklearn.metrics import confusion_matrix, roc_curve, auc
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.decomposition import PCA # For dimensionality reduction
(if necessary)
```

```
# Define features and target
```

```
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
```

```
# Example features
```

```
X = df[selected_features]
```

```
y = (df['Price (Euro)'] > df['Price (Euro)'].median()).astype(int) #
Binary classification
```

```
# Split data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

```
# Train SVM model
```

```
svm_clf = SVC(kernel='linear', probability=True, random_state=42) #
Change kernel to 'rbf', 'poly', etc., as needed
```

```
svm_clf.fit(X_train, y_train)
```

```


# Step 1: Confusion Matrix
y_pred = svm_clf.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
            xticklabels=["Below Median", "Above Median"],
            yticklabels=["Below Median", "Above Median"])
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()


# Step 2: Decision Boundary (if reduced to 2 features)
if len(selected_features) == 2:
    import numpy as np
    x_min, x_max = X_train[selected_features[0]].min() - 1,
    X_train[selected_features[0]].max() + 1
    y_min, y_max = X_train[selected_features[1]].min() - 1,
    X_train[selected_features[1]].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                        np.arange(y_min, y_max, 0.01))
    Z = svm_clf.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.figure(figsize=(8, 6))
    plt.contourf(xx, yy, Z, alpha=0.8, cmap="coolwarm")
    sns.scatterplot(x=X_train[selected_features[0]],
                    y=X_train[selected_features[1]],
                    hue=y_train, palette="coolwarm", edgecolor="k")
    plt.title("Decision Boundary (SVM)")
    plt.xlabel(selected_features[0])
    plt.ylabel(selected_features[1])
    plt.show()

# Step 3: ROC Curve
y_pred_proba = svm_clf.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC Curve (AUC =
{roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--')
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid()
plt.show()

```





```
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Step 1: Select the features for clustering
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
                    "CPU_Frequency (GHz)"] # Features to use for clustering
X = df[selected_features]

# Step 2: Standardize the features (K-means is sensitive to feature
# scaling)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Apply K-means clustering (let's assume 3 clusters, you can
# tune this)
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(X_scaled)

# Step 4: Add cluster labels to the dataframe
```

```

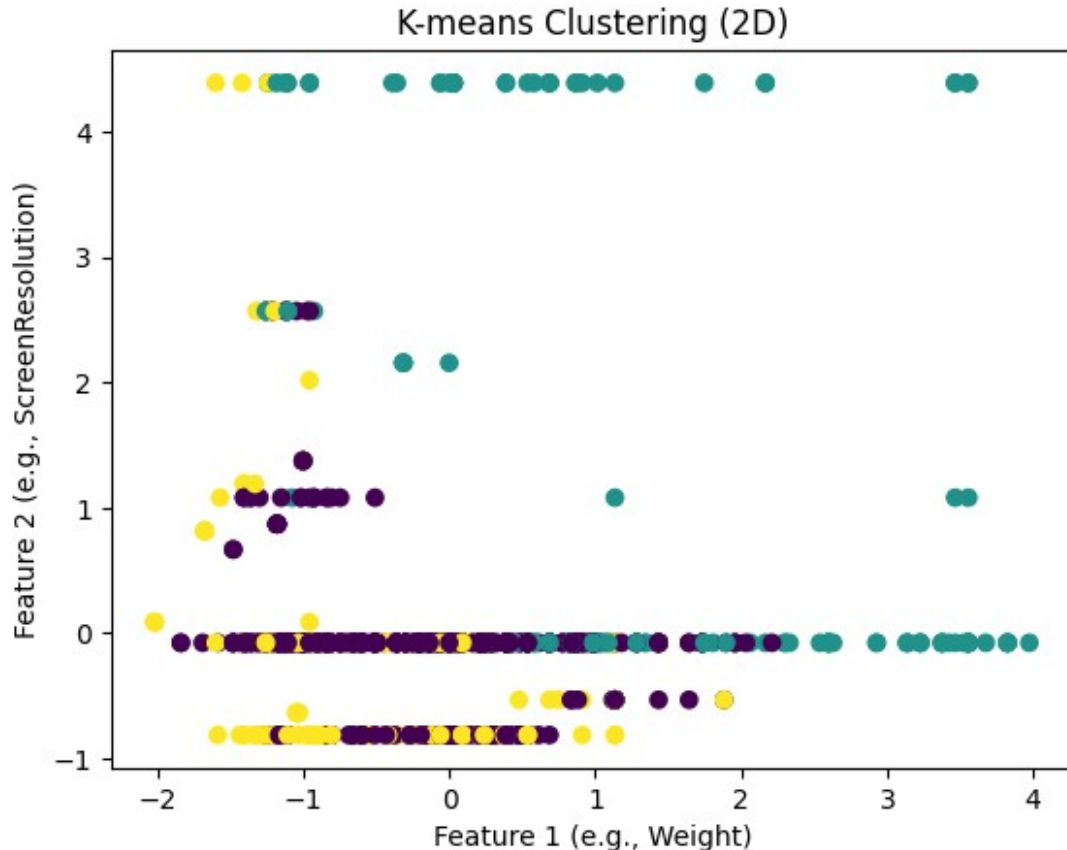
df['Cluster'] = kmeans.labels_

# Step 5: Evaluate clustering - inertia (sum of squared distances of
samples to their closest cluster center)
print(f"Inertia: {kmeans.inertia_:.4f}")

# Step 6: Optional - Visualize the clusters (2D visualization for
simplicity)
plt.scatter(X_scaled[:, 0], X_scaled[:, 1], c=kmeans.labels_,
cmap='viridis')
plt.title('K-means Clustering (2D)')
plt.xlabel('Feature 1 (e.g., Weight)')
plt.ylabel('Feature 2 (e.g., ScreenResolution)')
plt.show()

Inertia: 2891.4306

```



```

from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

```



```

# Define features for clustering
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
# Example features
X = df[selected_features]

# Step 1: Elbow Method for Optimal Number of Clusters
inertia = []
k_values = range(1, 11) # Try 1 to 10 clusters
for k in k_values:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X)
    inertia.append(kmeans.inertia_)

plt.figure(figsize=(8, 6))
plt.plot(k_values, inertia, marker='o', linestyle='--')
plt.title("Elbow Method for Optimal K")
plt.xlabel("Number of Clusters (K)")
plt.ylabel("Inertia")
plt.grid()
plt.show()

# Step 2: K-Means Clustering
optimal_k = 3 # Based on elbow method (adjust as needed)
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
clusters = kmeans.fit_predict(X)
df["Cluster"] = clusters # Add cluster labels to the dataset

# Step 3: Silhouette Score
sil_score = silhouette_score(X, clusters)
print(f"Silhouette Score: {sil_score:.2f}")

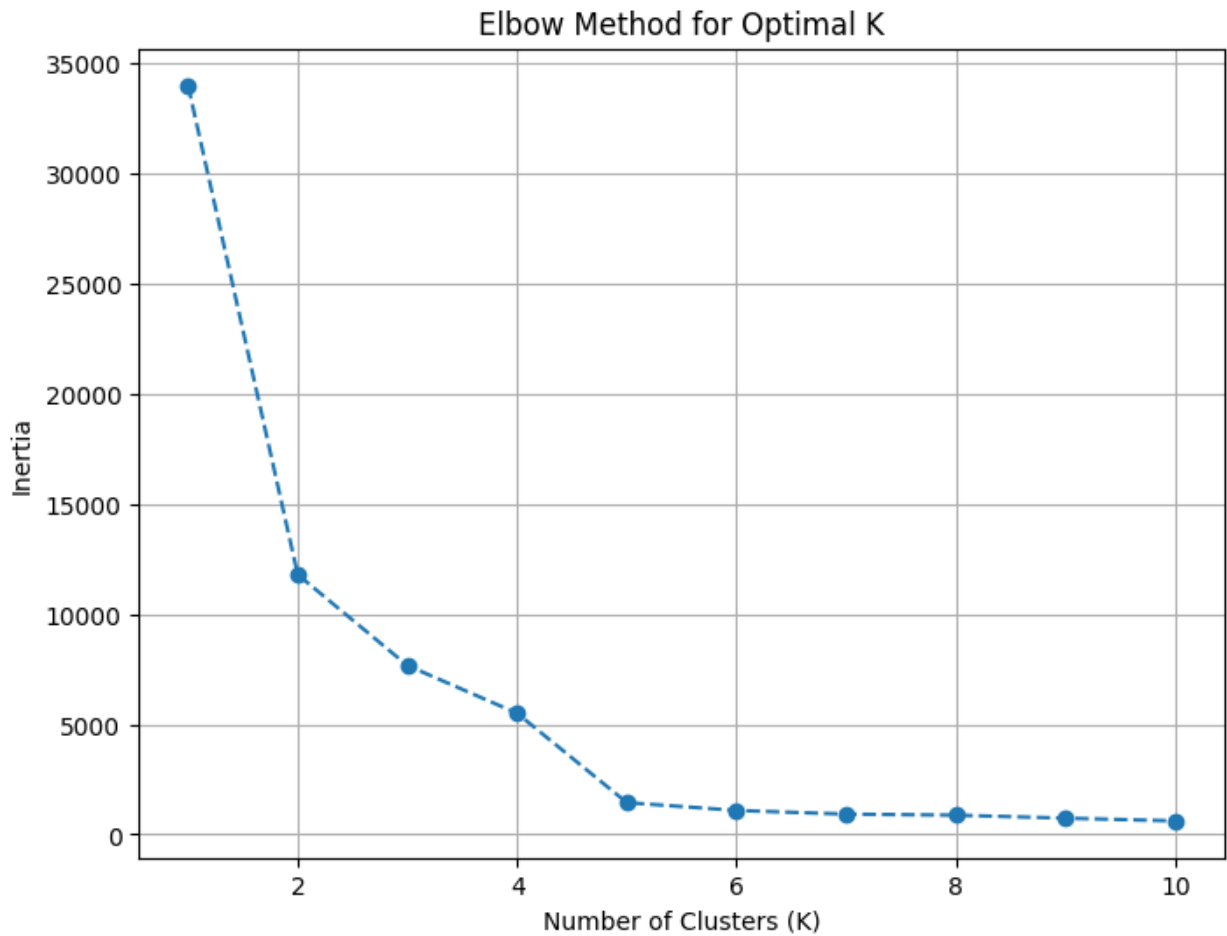
# Step 4: PCA for Visualization (Reduce to 2D)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X)

plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_pca[:, 0], y=X_pca[:, 1], hue=clusters,
palette="coolwarm", s=100, edgecolor="k")
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], s=300, c='gold', label='Centroids')
plt.title("K-Means Clustering (2D Visualization with PCA)")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.legend()
plt.grid()
plt.show()

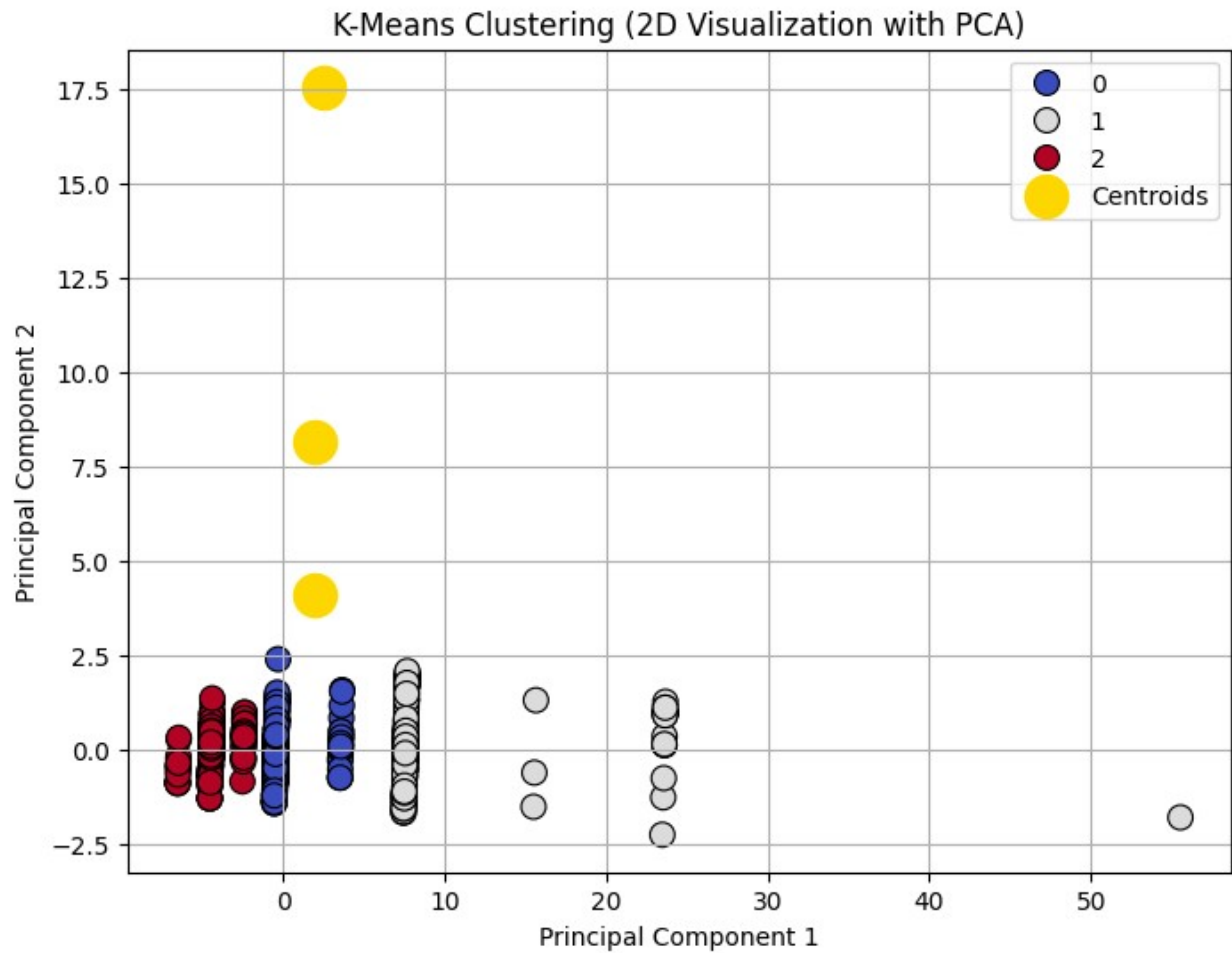
# Optional: Distribution of Clusters
plt.figure(figsize=(6, 4))

```

```
sns.countplot(x="Cluster", data=df, palette="coolwarm")  
plt.title("Cluster Distribution")  
plt.xlabel("Cluster")  
plt.ylabel("Number of Points")  
plt.show()
```




Silhouette Score: 0.71



```
<ipython-input-11-f60b85c4fe27>:54: FutureWarning:
```


```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.
```

```
sns.countplot(x="Cluster", data=df, palette="coolwarm")
```



```
inertia_values = []
for k in range(1, 11): # Testing clusters from 1 to 10
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_scaled)
    inertia_values.append(kmeans.inertia_)

# Plotting the elbow curve
plt.plot(range(1, 11), inertia_values, marker='o')
plt.title('Elbow Method for Optimal K')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.show()
```



```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report

# Step 1: Convert 'Price (Euro)' to binary classification (e.g., high
# price vs low price)
threshold = df['Price (Euro)'].median()
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features for KNN
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 4: Standardize the features (KNN is sensitive to feature
scaling)
```

```

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 5: Initialize the KNN classifier
knn_classifier = KNeighborsClassifier(n_neighbors=5) # You can tune
the number of neighbors (k)

# Step 6: Train the KNN model
knn_classifier.fit(X_train_scaled, y_train)

# Step 7: Make predictions on the test set
y_pred = knn_classifier.predict(X_test_scaled)

# Step 8: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.8706

Classification Report:

	precision	recall	f1-score	support
0	0.86	0.86	0.86	118
1	0.88	0.88	0.88	137
accuracy			0.87	255
macro avg	0.87	0.87	0.87	255
weighted avg	0.87	0.87	0.87	255

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, roc_curve, auc
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Step 1: Define features and target
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
# Example features
X = df[selected_features]
y = (df['Price (Euro)'] > df['Price (Euro)'].median()).astype(int) #
Binary classification: Above/Below Median Price

# Step 2: Standardize features for KNN
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

```

```

# Step 3: Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.3, random_state=42)

# Step 4: Train KNN model
k = 5 # You can tune this parameter
knn_clf = KNeighborsClassifier(n_neighbors=k)
knn_clf.fit(X_train, y_train)

# Step 5: Confusion Matrix Visualization
y_pred = knn_clf.predict(X_test)
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
xticklabels=["Below Median", "Above Median"],
yticklabels=["Below Median", "Above Median"])
plt.title(f"Confusion Matrix (KNN, k={k})")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

# Step 6: ROC Curve
y_pred_proba = knn_clf.predict_proba(X_test)[:, 1] # Probability for
the positive class
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC Curve (AUC =
{roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='grey', linestyle='--')
plt.title("ROC Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.grid()
plt.show()

# Step 7: Decision Boundary (if reduced to 2D)
if len(selected_features) > 2: # Reduce features to 2D for
visualization
    pca = PCA(n_components=2)
    X_train_2d = pca.fit_transform(X_train)
    X_test_2d = pca.transform(X_test)
    knn_clf.fit(X_train_2d, y_train)
    x_min, x_max = X_train_2d[:, 0].min() - 1, X_train_2d[:, 0].max()
+ 1
    y_min, y_max = X_train_2d[:, 1].min() - 1, X_train_2d[:, 1].max()
+ 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),

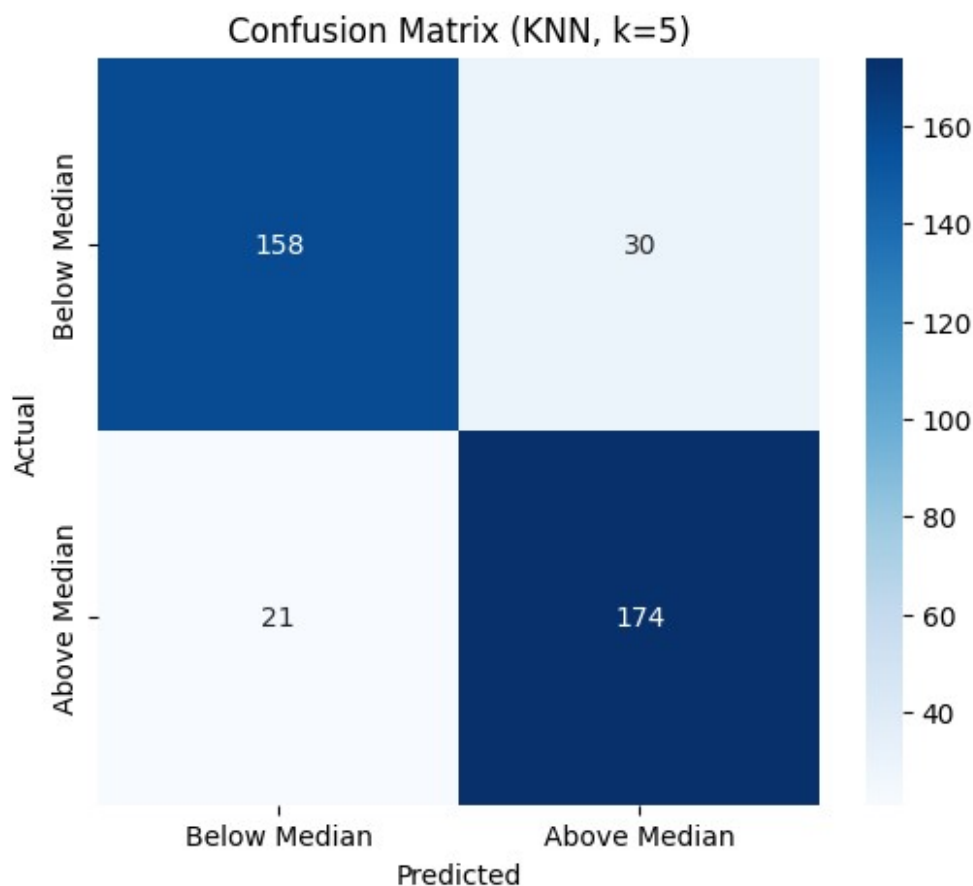
```

```


np.arange(y_min, y_max, 0.01)) # Fixed
closing parenthesis here
Z = knn_clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)


plt.figure(figsize=(8, 6))
plt.contourf(xx, yy, Z, alpha=0.8, cmap="coolwarm")
sns.scatterplot(x=X_train_2d[:, 0], y=X_train_2d[:, 1],
hue=y_train, palette="coolwarm", edgecolor="k")
plt.title("Decision Boundary (KNN)")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.show()

```









```

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report

# Step 1: Convert 'Price (Euro)' to binary classification (e.g., high
# price vs low price)
threshold = df['Price (Euro)'].median()
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features for LDA
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

```
# Step 4: Standardize the features (LDA is sensitive to feature scaling)
```

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Step 5: Initialize the LDA model
```

```
lda_classifier = LinearDiscriminantAnalysis()
```

```
# Step 6: Train the LDA model
```

```
lda_classifier.fit(X_train_scaled, y_train)
```

```
# Step 7: Make predictions on the test set
```

```
y_pred = lda_classifier.predict(X_test_scaled)
```

```
# Step 8: Evaluate the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy:.4f}")
```

```
print("Classification Report:")
```

```
print(classification_report(y_test, y_pred))
```

Accuracy: 0.7765

Classification Report:

	precision	recall	f1-score	support
0	0.73	0.81	0.77	118
1	0.82	0.74	0.78	137
accuracy			0.78	255
macro avg	0.78	0.78	0.78	255
weighted avg	0.78	0.78	0.78	255

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import numpy as np
```

```
# Step 1: Define features and target
```

```
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
```

```
# Example features
```

```
X = df[selected_features]
```

```
y = (df['Price (Euro)'] > df['Price (Euro)'].median()).astype(int) #
```

```
Binary classification: Above/Below Median Price
```

```
# Step 2: Standardize features for LDA
```

```
scaler = StandardScaler()
```

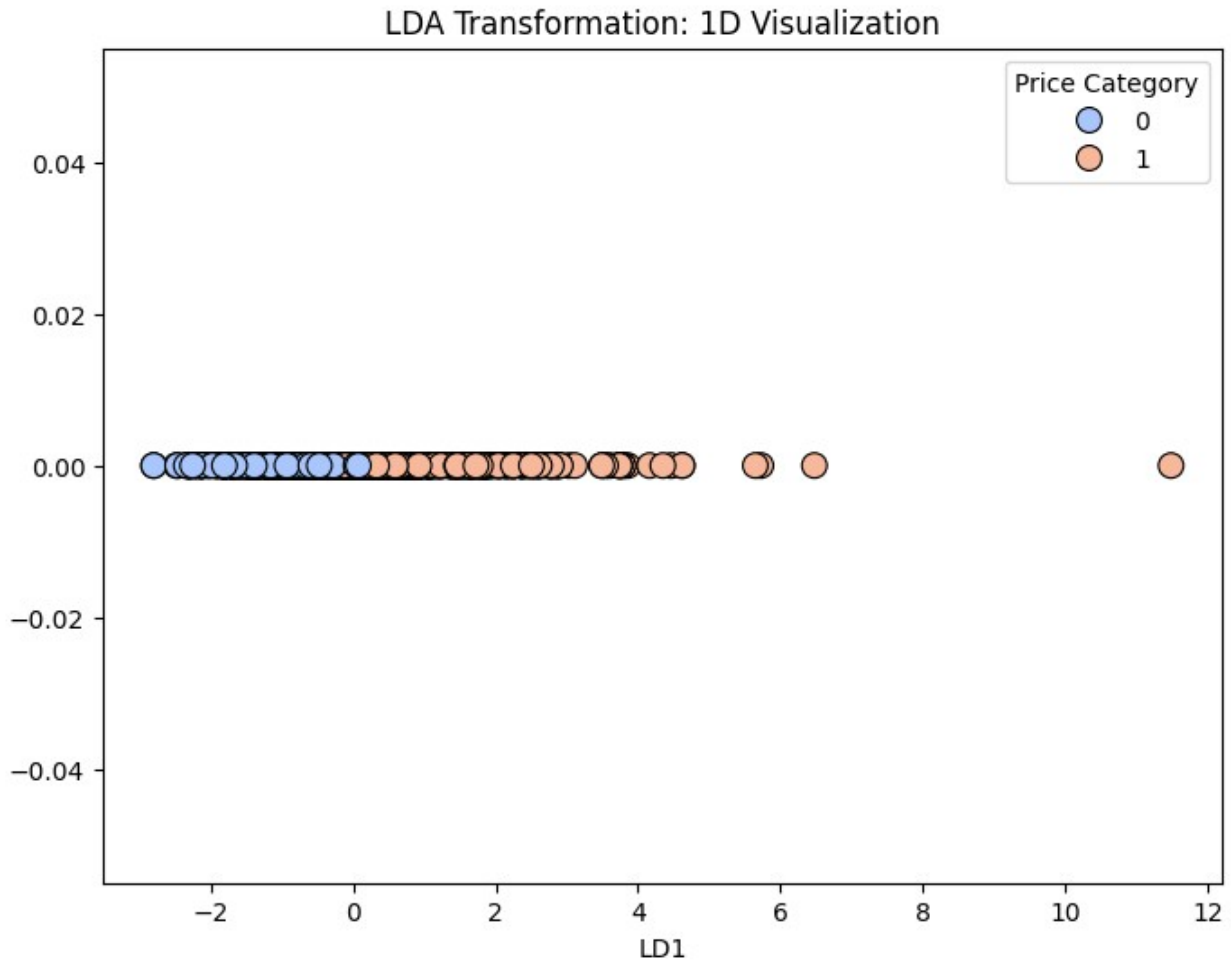
```
X_scaled = scaler.fit_transform(X)
```

```
# Step 3: Apply LDA with n_components=1
lda = LinearDiscriminantAnalysis(n_components=1) # Reduce to 1D since
we have only 2 classes
X_lda = lda.fit_transform(X_scaled, y)

# Step 4: Split data into training and testing sets (not necessary for
visualization but common practice)
X_train, X_test, y_train, y_test = train_test_split(X_lda, y,
test_size=0.3, random_state=42)

# Step 5: Visualization of LDA transformation in 1D (since only 1
component is chosen)
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_lda[:, 0], y=np.zeros_like(X_lda[:, 0]), hue=y,
palette='coolwarm', edgecolor='k', s=100)
plt.title("LDA Transformation: 1D Visualization")
plt.xlabel('LD1') # First Linear Discriminant Component
plt.ylabel('')
plt.legend(title="Price Category", loc='best')
plt.show()

# Optionally, visualize decision boundaries for LDA if needed (for 1D
it may not be meaningful)
```



```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report
from sklearn.tree import DecisionTreeClassifier

# Step 1: Convert 'Price (Euro)' to binary classification (e.g., high
# price vs low price)
threshold = df['Price (Euro)'].median() # Define the threshold for
classification
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features for AdaBoost
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Split the data into training and testing sets
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 4: Standardize the features (AdaBoost works better when
features are scaled)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 5: Initialize the AdaBoost model (using default decision stump
base estimator)
ada_boost_classifier = AdaBoostClassifier(n_estimators=50,
random_state=42)

# Step 6: Train the AdaBoost model
ada_boost_classifier.fit(X_train_scaled, y_train)

# Step 7: Make predictions on the test set
y_pred = ada_boost_classifier.predict(X_test_scaled)

# Step 8: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/
_weight_boosting.py:527: FutureWarning: The SAMME.R algorithm (the
default) is deprecated and will be removed in 1.6. Use the SAMME
algorithm to circumvent this warning.

```

```
warnings.warn(
```

Accuracy: 0.8510

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.85	0.84	118
1	0.87	0.85	0.86	137
accuracy			0.85	255
macro avg	0.85	0.85	0.85	255
weighted avg	0.85	0.85	0.85	255

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, roc_curve, auc,
confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.ensemble import AdaBoostClassifier

```

```

import pandas as pd

# Step 1: Convert target variable 'Price (Euro)' to binary
classification (e.g., high price vs low price)
threshold = df['Price (Euro)'].median() # Define a threshold for
binary classification (e.g., median price)
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features and handle the categorical
'ScreenResolution' feature
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Convert 'ScreenResolution' (categorical) to numerical format
using LabelEncoder
label_encoder = LabelEncoder()
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 5: Standardize the features (optional but recommended for
AdaBoost)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 6: Initialize AdaBoost Classifier (no base estimator specified)
ada_boost_classifier = AdaBoostClassifier(n_estimators=50,
random_state=42)

# Step 7: Train the model
ada_boost_classifier.fit(X_train_scaled, y_train)

# Step 8: Make predictions on the test set
y_pred = ada_boost_classifier.predict(X_test_scaled)
y_pred_prob = ada_boost_classifier.predict_proba(X_test_scaled)[: , 1]
# Probabilities for the positive class

# Step 9: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Visualizations

```

```

## 1. Accuracy of AdaBoost Classifier
print("Accuracy of AdaBoost Classifier: ", accuracy)

## 2. ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC Curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='red', linestyle='--') # Diagonal line (Random classifier)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

## 3. Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Low Price', 'High Price'], yticklabels=['Low Price', 'High Price'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for AdaBoost Classifier')
plt.show()

```

<ipython-input-28-0e8ca5f63704>:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation:  
[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_weight_boost
ing.py:527: FutureWarning: The SAMME.R algorithm (the default) is
deprecated and will be removed in 1.6. Use the SAMME algorithm to
circumvent this warning.
warnings.warn(


```


Accuracy: 0.8549

Accuracy of AdaBoost Classifier: 0.8549019607843137



Receiver Operating Characteristic (ROC) Curve





```
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report

# Step 1: Convert 'Price (Euro)' to binary classification (e.g., high
# price vs low price)
threshold = df['Price (Euro)'].median() # Define the threshold for
classification
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features for the model
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Example features
X = df[selected_features] # Features
y = y_binary # Target (binary classification)
```

```

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 4: Standardize the features (though Gradient Boosting is less
sensitive to this)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 5: Initialize the Gradient Boosting model
gradient_boosting_classifier =
GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
random_state=42)

# Step 6: Train the Gradient Boosting model
gradient_boosting_classifier.fit(X_train_scaled, y_train)

# Step 7: Make predictions on the test set
y_pred = gradient_boosting_classifier.predict(X_test_scaled)

# Step 8: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.8588

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.88	0.85	118
1	0.89	0.84	0.86	137
accuracy			0.86	255
macro avg	0.86	0.86	0.86	255
weighted avg	0.86	0.86	0.86	255

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, roc_curve, auc
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import pandas as pd

```

```

# Step 1: Convert target variable 'Price (Euro)' to binary
classification (e.g., high price vs low price)
threshold = df['Price (Euro)'].median() # Define a threshold for
binary classification (e.g., median price)

```

```

y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features and handle the categorical
'ScreenResolution' feature
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Convert 'ScreenResolution' (categorical) to numerical format
using LabelEncoder
label_encoder = LabelEncoder()
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 5: Standardize the features (optional)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 6: Initialize Gradient Boosting Classifier
gb_classifier = GradientBoostingClassifier(n_estimators=100,
random_state=42)

# Step 7: Train the model
gb_classifier.fit(X_train_scaled, y_train)

# Step 8: Make predictions on the test set
y_pred = gb_classifier.predict(X_test_scaled)
y_pred_prob = gb_classifier.predict_proba(X_test_scaled)[: , 1] #
Probabilities for the positive class

# Step 9: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Visualizations

## 1. Accuracy of Gradient Boosting Classifier
print("Accuracy of Gradient Boosting Classifier: ", accuracy)

## 2. ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)

```

```
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC Curve (area =
{roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='red', linestyle='--') # Diagonal line
(Random classifier)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()
```

### ## 3. Feature Importance Plot

```
importances = gb_classifier.feature_importances_
feature_names = X.columns
```

```
plt.figure(figsize=(8, 6))
sns.barplot(x=importances, y=feature_names)
plt.title('Feature Importance for Gradient Boosting Classifier')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```

<ipython-input-25-9440e78996d2>:20: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation:


[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

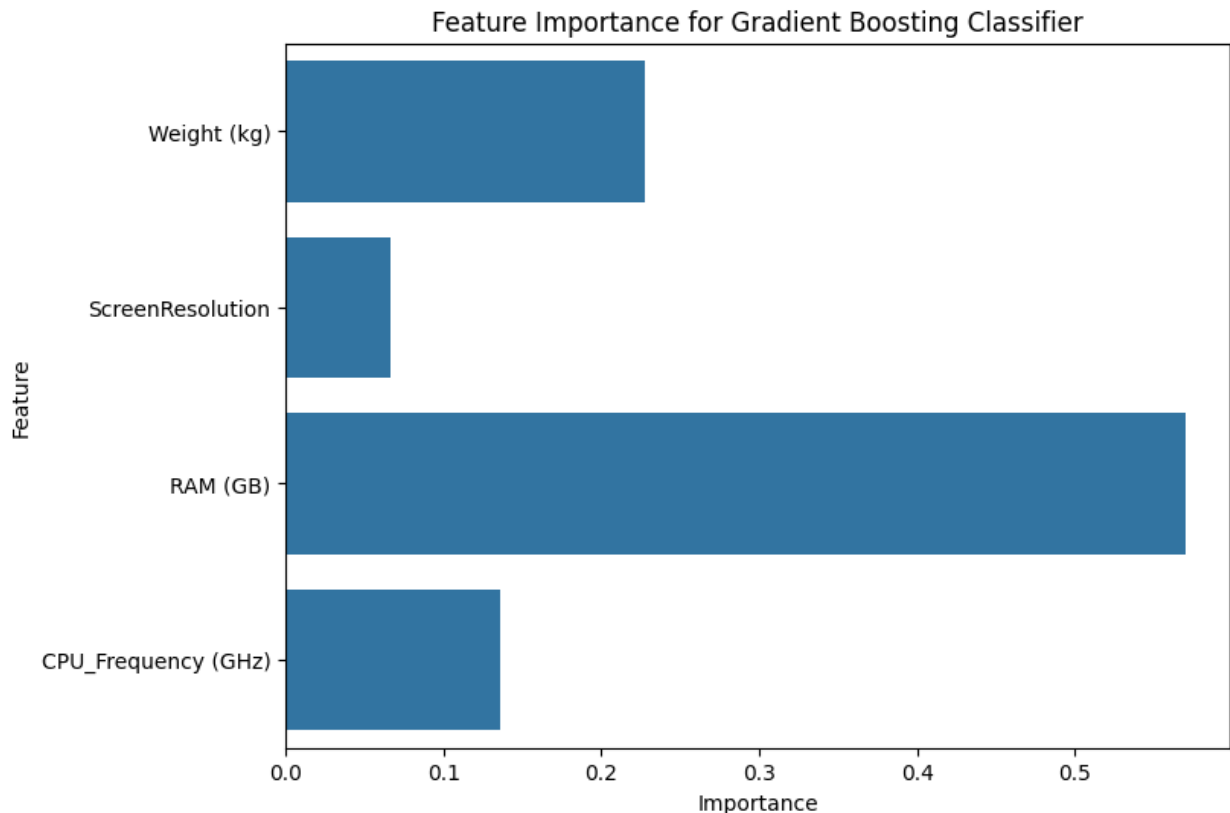
```
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))
```

Accuracy: 0.8588

Accuracy of Gradient Boosting Classifier: 0.8588235294117647

Receiver Operating Characteristic (ROC) Curve





```
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report

# Step 1: Convert 'Price (Euro)' to binary classification (e.g., high
price vs low price)
threshold = df['Price (Euro)'].median() # Define the threshold for
classification
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features for the model
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Example features
X = df[selected_features] # Features
y = y_binary # Target (binary classification)

# Step 3: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 4: Standardize the features (Gaussian Naive Bayes assumes
normally distributed features)
scaler = StandardScaler()
```

```

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 5: Initialize the Naive Bayes model
naive_bayes_classifier = GaussianNB()

# Step 6: Train the Naive Bayes model
naive_bayes_classifier.fit(X_train_scaled, y_train)

# Step 7: Make predictions on the test set
y_pred = naive_bayes_classifier.predict(X_test_scaled)

# Step 8: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")
print("Classification Report:")
print(classification_report(y_test, y_pred))

```

Accuracy: 0.7098

Classification Report:

	precision	recall	f1-score	support
0	0.62	0.97	0.75	118
1	0.94	0.49	0.64	137
accuracy			0.71	255
macro avg	0.78	0.73	0.70	255
weighted avg	0.79	0.71	0.70	255

```

import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import accuracy_score, roc_curve, auc,
confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.naive_bayes import GaussianNB
import pandas as pd

# Step 1: Convert target variable 'Price (Euro)' to binary
classification (e.g., high price vs low price)
threshold = df['Price (Euro)'].median() # Define a threshold for
binary classification (e.g., median price)
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price
> threshold, 0 if price <= threshold

# Step 2: Select features and handle the categorical
'ScreenResolution' feature
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",
"CPU_Frequency (GHz)"] # Features
X = df[selected_features] # Features

```



```

y = y_binary # Target (binary classification)

# Step 3: Convert 'ScreenResolution' (categorical) to numerical format
using LabelEncoder
label_encoder = LabelEncoder()
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))

# Step 4: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Step 5: Standardize the features (optional but recommended for Naive
Bayes)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 6: Initialize Naive Bayes Classifier
nb_classifier = GaussianNB()

# Step 7: Train the model
nb_classifier.fit(X_train_scaled, y_train)

# Step 8: Make predictions on the test set
y_pred = nb_classifier.predict(X_test_scaled)
y_pred_prob = nb_classifier.predict_proba(X_test_scaled)[: , 1] #
Probabilities for the positive class

# Step 9: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Visualizations

## 1. Accuracy of Naive Bayes Classifier
print("Accuracy of Naive Bayes Classifier: ", accuracy)

## 2. ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='blue', label=f'ROC Curve (area =
{roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='red', linestyle='--') # Diagonal line
(Random classifier)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')

```

```
plt.legend(loc='lower right')
plt.show()
```

### *## 3. Confusion Matrix*

```
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Low Price', 'High Price'], yticklabels=['Low Price',
            'High Price'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Naive Bayes Classifier')
plt.show()
```

```
<ipython-input-26-2df5ca840585>:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:


[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))
```

Accuracy: 0.7529

Accuracy of Naive Bayes Classifier: 0.7529411764705882

Receiver Operating Characteristic (ROC) Curve





```
from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.naive_bayes import GaussianNB

# Step 1: Define features and target
selected_features = ["Weight (kg)", "RAM (GB)", "CPU_Frequency (GHz)"]
# Example features
X = df[selected_features]
y = (df['Price (Euro)'] > df['Price (Euro)'].median()).astype(int) #
```

*Binary classification: Above/Below Median Price*

*# Step 2: Standardize features*

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

*# Step 3: Split the data into training and testing sets*

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.3, random_state=42)
```

*# Step 4: Initialize models*

```
models = {
    'Logistic Regression': LogisticRegression(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'SVM': SVC(),
    'KNN': KNeighborsClassifier(),
    'LDA': LinearDiscriminantAnalysis(),
    'AdaBoost': AdaBoostClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'Naive Bayes': GaussianNB()
}
```

*# Step 5: Train and evaluate each model*

```
results = {}
for model_name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # Calculate performance metrics
    accuracy = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)

    results[model_name] = {
        'Accuracy': accuracy,
        'F1 Score': f1,
        'Precision': precision,
        'Recall': recall
    }
```

*# Step 6: Display results*

```
results_df = pd.DataFrame(results).T
print(results_df.sort_values(by='F1 Score', ascending=False))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/
_weight_boosting.py:527: FutureWarning: The SAMME.R algorithm (the
default) is deprecated and will be removed in 1.6. Use the SAMME
```

```
algorithm to circumvent this warning.  
warnings.warn(  

```

	Accuracy	F1 Score	Precision	Recall
AdaBoost	0.877285	0.877922	0.889474	0.866667
KNN	0.866841	0.872180	0.852941	0.892308
Random Forest	0.861619	0.862338	0.873684	0.851282
SVM	0.859008	0.862245	0.857868	0.866667
Gradient Boosting	0.856397	0.857143	0.868421	0.846154
Decision Tree	0.838120	0.836842	0.859459	0.815385
Logistic Regression	0.801567	0.814634	0.776744	0.856410
LDA	0.798956	0.798956	0.813830	0.784615
Naive Bayes	0.767624	0.722741	0.920635	0.594872

```
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.ensemble import RandomForestClassifier,  
AdaBoostClassifier, GradientBoostingClassifier  
from sklearn.svm import SVC  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.linear_model import LogisticRegression  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.naive_bayes import GaussianNB  
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis  
from sklearn.metrics import accuracy_score  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler, LabelEncoder  
  
# Step 1: Convert target variable 'Price (Euro)' to binary  
classification  
threshold = df['Price (Euro)'].median() # Define a threshold for  
binary classification  
y_binary = (df['Price (Euro)'] > threshold).astype(int) # 1 if price  
> threshold, 0 if price <= threshold  
  
# Step 2: Select features and handle the categorical  
'ScreenResolution' feature  
selected_features = ["Weight (kg)", "ScreenResolution", "RAM (GB)",  
"CPU_Frequency (GHz)"] # Features  
X = df[selected_features] # Features  
y = y_binary # Target (binary classification)  
  
# Convert 'ScreenResolution' (categorical) to numerical format using  
LabelEncoder  
label_encoder = LabelEncoder()  
X['ScreenResolution'] =  
label_encoder.fit_transform(X['ScreenResolution'].astype(str))  
  
# Step 3: Split the data into training and testing sets  
X_train, X_test, y_train, y_test = train_test_split(X, y,
```

```

test_size=0.2, random_state=42)

# Step 4: Standardize the features (optional but recommended for some
models)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Step 5: Initialize all classifiers
models = {
    'Logistic Regression': LogisticRegression(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(n_estimators=100,
random_state=42),
    'SVM': SVC(probability=True),
    'KNN': KNeighborsClassifier(),
    'Naive Bayes': GaussianNB(),
    'AdaBoost': AdaBoostClassifier(n_estimators=50, random_state=42),
    'Gradient Boosting': GradientBoostingClassifier(),
    'LDA': LinearDiscriminantAnalysis(),
}

# Step 6: Train each model, make predictions, and store accuracy
scores
accuracy_scores = []
model_names = []

for model_name, model in models.items():
    # Train model
    model.fit(X_train_scaled, y_train)
    # Make predictions
    y_pred = model.predict(X_test_scaled)
    # Calculate accuracy score
    accuracy = accuracy_score(y_test, y_pred)
    accuracy_scores.append(accuracy)
    model_names.append(model_name)

# Step 7: Create a DataFrame for visualization
model_comparison = pd.DataFrame({
    'Model': model_names,
    'Accuracy': accuracy_scores
})

# Step 8: Visualize the model comparison (Bar chart)
plt.figure(figsize=(10, 6))
sns.barplot(x='Accuracy', y='Model', data=model_comparison,
palette='viridis')
plt.title('Comparison of Model Performance')
plt.xlabel('Accuracy')

```

```
plt.ylabel('Model')
plt.show()
```

```
<ipython-input-29-450d9891bf53>:25: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X['ScreenResolution'] =
label_encoder.fit_transform(X['ScreenResolution'].astype(str))
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_weight_boost
ing.py:527: FutureWarning: The SAMME.R algorithm (the default) is
deprecated and will be removed in 1.6. Use the SAMME algorithm to
circumvent this warning.
```

```
warnings.warn(
<ipython-input-29-450d9891bf53>:70: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Accuracy', y='Model', data=model_comparison,
palette='viridis')
```

